

## QuantumLedger System Overview

**QuantumLedger** is a Java-based immutable ledger designed for accounting transactions. It consists of modular components built on Java SE with cryptographic hashing (SHA-256), using JDBC with an embedded SQLite database for persistence. The system's core modules include a **Block** (encapsulating timestamp, data, previous hash, and hash), a **Blockchain** (a linked list of blocks forming the ledger), a **LedgerService** for managing transactions (financial, inventory, general), and a **UserAuth** module for role-based access. The project is organized as a Maven project with a standard directory structure (e.g., `src/main/java/...` for source, `src/main/resources` for configuration, and `pom.xml` for dependencies).

### Maven Project Structure

A typical Maven layout for QuantumLedger might be:

```
quantumledger/
├─ pom.xml
├─ src/
│   └─ main/
│       ├── java/com/quantumledger/
│       │   ├── block/           # Block class, hashing logic
│       │   ├── blockchain/      # Blockchain class (LinkedList)
│       │   ├── service/         # LedgerService, DAOs
│       │   ├── auth/            # UserAuth, role management
│       │   └─ model/            # JPA entity classes (Block, Transaction, User,
│                               Role)
│       └─ resources/           # SQLite schema (DDL), config files
└─ test/java/                  # Unit tests
```

The **pom.xml** should include dependencies for the SQLite JDBC driver (e.g. [org.xerial:sqlite-jdbc](https://mvnrepository.com/artifact/org.xerial/sqlite-jdbc)) and optional JPA/Hibernate (for future ORM). For example:

```
<dependencies>
  <dependency>
    <groupId>org.xerial</groupId>
    <artifactId>sqlite-jdbc</artifactId>
    <version>3.43.0.0</version>
  </dependency>
  <!-- (Optional) Hibernate and JPA if migrating from JDBC -->
  <dependency>
    <groupId>org.hibernate.orm</groupId>
    <artifactId>hibernate-community-dialects</artifactId>
    <version>6.4.1.Final</version>
  </dependency>
</dependencies>
```

```
</dependency>
</dependencies>
```

These provide the embedded SQLite driver and, for Hibernate 6.x, a built-in SQLite dialect <sup>1</sup>.

## Block Module

The **Block** class defines the structure of a ledger block. It includes fields like: - `private String data;` - the payload (e.g. transaction details). - `private long timeStamp;` - milliseconds since epoch. - `public String previousHash;` - the hash of the prior block. - `public String hash;` - the SHA-256 hash of this block.

On creation, the `Block` computes its hash by concatenating its data, timestamp, and previous hash, then hashing with SHA-256. For example (inspired by standard implementations) <sup>2</sup>:

```
public class Block {
    public String hash;
    public String previousHash;
    private String data;
    private long timeStamp;

    public Block(String data, String previousHash) {
        this.data = data;
        this.previousHash = previousHash;
        this.timeStamp = System.currentTimeMillis();
        this.hash = calculateHash();
    }

    // Calculate SHA-256 hash of (previousHash + timestamp + data)
    public String calculateHash() {
        // Example pseudo-code for hashing
        MessageDigest sha = MessageDigest.getInstance("SHA-256");
        String input = previousHash + Long.toString(timeStamp) + data;
        byte[] hashedBytes = sha.digest(input.getBytes("UTF-8"));
        return bytesToHex(hashedBytes);
    }
}
```

(Code simplified for brevity.) The use of `MessageDigest.getInstance("SHA-256")` is standard for computing SHA-256 in Java <sup>3</sup>. The `calculateHash()` method returns a hexadecimal string of the hash. Changing any field (data, timestamp, or previousHash) would change the block's hash, ensuring integrity.

## Blockchain Module

The **Blockchain** class maintains the ledger as a chain of `Block` objects. Internally it can use a `LinkedList<Block>` (or `ArrayList<Block>`) to store the blocks in order <sup>4</sup>. For example:

```

public class Blockchain {
    private List<Block> chain;

    public Blockchain() {
        this.chain = new LinkedList<>(); // tamper-evident chain
    }

    public void addBlock(Block block) {
        chain.add(block);
    }

    public List<Block> getChain() {
        return chain;
    }
}

```

Using a `LinkedList` reflects the natural “chain” structure <sup>4</sup>. Each new block’s `previousHash` links to the hash of the last block in the chain, forming an immutable linked list. This design means that once a block is added, any alteration breaks the hashes: in a correct chain, each block’s stored hash should equal its recalculated hash, and each block’s `previousHash` should match the prior block’s actual hash. Violations of these conditions are easy to detect, making the ledger *tamper-evident*. As one source explains, blockchain data is “resistant to modification” – once a block is added, its data cannot be changed without invalidating subsequent blocks <sup>5</sup> <sup>6</sup>.

## LedgerService (Transactions)

The **LedgerService** provides operations to add and view transactions. Each transaction has a type (e.g. `FINANCIAL`, `INVENTORY`, `GENERAL`) and associated details (e.g. amount, items). When adding a transaction, the service: 1. Creates a `Block` containing the transaction data (e.g. as a JSON or concatenated string). 2. Appends it to the `Blockchain` (managing the linked-list in memory). 3. Persists the transaction (and block metadata) to the SQLite database via JDBC.

The service should validate role permissions (from `UserAuth`) and then use a Data Access Object (DAO) layer for database writes <sup>7</sup>. The DAO pattern keeps SQL operations separated from business logic <sup>7</sup>. For example:

```

// Pseudo-code in LedgerService
public void recordTransaction(Transaction txn, User user) {
    if (!user.hasRole("Accountant")) throw new AuthException();
    Block newBlock = new Block(txn.toString(), getLatestBlockHash());
    blockchain.addBlock(newBlock);
    transactionDao.save(txn, newBlock.getHash()); // DAO writes to SQLite
}

```

Optional soft reversibility is supported by *adding* compensating entries instead of deleting history. In a double-entry ledger style, one might **reverse** an erroneous transaction by creating a new reversing transaction (e.g. negative of the original) rather than deleting the original block. As explained by accounting systems design, if an error occurs the “ledger transaction stays immutable” and you either create a new reverse transaction with the opposite amount or a delta transaction to correct it <sup>8</sup>. Thus

QuantumLedger would never delete or alter an existing block; it only appends new blocks to reflect reversals or adjustments, ensuring a full audit trail.

## UserAuth Module

The **UserAuth** module manages users and roles (e.g. `Admin`, `Accountant`, `Auditor`). A normalized schema typically uses separate **Users** and **Roles** tables, with a join table (e.g. `user_roles`) for many-to-many relationships. For simplicity, one can start with each user having one role. For example:

- `users(id PK, username, password_hash, email)`
- `roles(id PK, name)`
- `user_roles(user_id FK → users.id, role_id FK → roles.id)`

Passwords should be stored as hashes (e.g. using PBKDF2 or bcrypt) rather than plaintext. This module checks credentials on login and associates each authenticated session with a user role. The `LedgerService` then restricts actions based on these roles (e.g. only `Admin` or `Accountant` can add transactions).

## Database Schema (SQLite)

A normalized SQLite schema for QuantumLedger might include tables for blocks, transactions, users, roles, and their relationships. Example tables:

- **blocks:** `block_id (PK)`, `timestamp`, `data`, `previous_hash`, `hash`.
- **transactions:** `txn_id (PK)`, `block_id (FK → blocks.block_id)`, `type`, `amount`, `description`, `txn_timestamp`.
- **users:** `user_id (PK)`, `username (unique)`, `password_hash`, `email`.
- **roles:** `role_id (PK)`, `role_name`.
- **user\_roles:** `user_id (FK)`, `role_id (FK)`, (composite PK).

Foreign keys enforce integrity (e.g. each transaction links to a block, each user-role entry links to valid user and role). This design avoids “storing two different entities in one table,” which is a common normalization error <sup>9</sup>. For example, one should *not* mix user credentials and transaction data in the same table. Each table is normalized to avoid redundancy and maintain consistency.

The entity-relationship structure is: **Block 1-N Transaction** (each block may contain one or more transactions), and **User M-N Role** via `user_roles`. (See ER diagram below for illustration.) By design, once a row is inserted in, say, `blocks` or `transactions`, it is never deleted; “archiving” a transaction would be done by appending a new block indicating reversal, rather than removing the old one.

## Example SQLite DDL

```
CREATE TABLE roles (  
    role_id    INTEGER PRIMARY KEY,  
    role_name TEXT NOT NULL UNIQUE  
);  
  
CREATE TABLE users (  
    user_id    INTEGER PRIMARY KEY,  
    username   TEXT NOT NULL UNIQUE,  
    password_hash TEXT NOT NULL,  
    email      TEXT NOT NULL UNIQUE  
);
```

```

    user_id      INTEGER PRIMARY KEY,
    username     TEXT NOT NULL UNIQUE,
    password_hash TEXT NOT NULL,
    email        TEXT
);

CREATE TABLE user_roles (
    user_id INTEGER,
    role_id INTEGER,
    PRIMARY KEY (user_id, role_id),
    FOREIGN KEY (user_id) REFERENCES users(user_id),
    FOREIGN KEY (role_id) REFERENCES roles(role_id)
);

CREATE TABLE blocks (
    block_id      INTEGER PRIMARY KEY AUTOINCREMENT,
    timestamp     INTEGER NOT NULL,
    data          TEXT,
    previous_hash TEXT,
    hash          TEXT NOT NULL
);

CREATE TABLE transactions (
    txn_id        INTEGER PRIMARY KEY AUTOINCREMENT,
    block_id      INTEGER,
    type          TEXT,
    amount        REAL,
    description    TEXT,
    txn_timestamp INTEGER NOT NULL,
    FOREIGN KEY (block_id) REFERENCES blocks(block_id)
);

```

(An ER diagram for this schema would show Block linked to Transaction, and User linked to Role via User\_Role.)

## Security and Immutability

QuantumLedger ensures data integrity via the cryptographic chain of hashes. Once a block is added, **its data is immutable** – changing it would require rehashing it and all subsequent blocks, which is easily detected <sup>5</sup>. In practice, “the data within [a block] is considered unalterable” <sup>6</sup>. To handle corrections or reversals, the system uses *append-only* compensating entries: e.g. a new transaction/block is created with negative amounts to offset an erroneous entry, leaving the original block intact. As noted in financial ledger design, immutability enforces an audit trail: even if mistakes occur, a new transaction is added to *reverse* the mistake rather than deleting history <sup>8</sup> <sup>6</sup>. This creates a “paper trail” of all actions, which is essential for auditability and trust.

## Multi-Tier Frontend Support

QuantumLedger's design allows multiple frontends:

- **Desktop GUI (JavaFX or Swing):** A standalone Java application can use the core libraries directly. The UI can call `LedgerService` methods, which update the in-memory blockchain and SQLite DB.
- **Web Backend (REST API):** Expose the ledger operations over HTTP. Using JAX-RS or Spring Boot, define endpoints like `POST /transactions` or `GET /blocks`. The backend calls the same service layer. For example, a REST controller might use `@Path("/transactions")` and call `ledgerService.addTransaction(...)`.
- **Mobile Integration:** Mobile apps (Android/iOS) can interact with a REST API or local database. For instance, an Android app could use HTTP to the Java backend, or even embed the same SQLite file and use Java code via JNI or Kotlin for core logic. The key is that all application types share the same data model and service interface, ensuring consistency.

The decentralized core allows each client to run independently or connect to a central server. For example, the SQLite database file can be synchronized or shared via networked file storage if offline modes are needed.

## JDBC and JPA Mapping

Currently, QuantumLedger uses plain JDBC to interact with SQLite (via the Xerial driver). SQL queries or simple DAO classes perform CRUD operations on the tables. In the future, one could migrate to JPA/Hibernate for ORM benefits. JPA entities would be annotated with `@Entity` and mapped to the same tables. For example:

```
@Entity @Table(name="blocks")
public class BlockEntity {
    @Id @GeneratedValue
    private Long blockId;
    private Long timestamp;
    private String data;
    private String previousHash;
    private String hash;
    // getters/setters...
}
```

Similarly for `TransactionEntity`, `UserEntity`, `RoleEntity`. Using JPA simplifies queries and caching. As of Hibernate 6, SQLite is officially supported via a community dialect (configure `hibernate.dialect = org.hibernate.community.dialect.SQLiteDialect`) <sup>1</sup>. With this setup, Hibernate can use `@ManyToOne` for `block` → `transactions` or `@ManyToMany` for `users` → `roles`. The database schema remains the same; only the data access layer is refactored. Migrating to JPA allows using the same entity classes across JDBC and ORM, easing future development.

## Design Patterns and Best Practices

QuantumLedger employs several established patterns:

- **Singleton (Blockchain):** The ledger chain itself can be managed as a singleton to ensure a single in-memory instance of `Blockchain` is used application-wide <sup>10</sup>. This prevents multiple conflicting chains.
- **Factory (Block creation):** A factory or builder could abstract block creation. For example, a `BlockFactory` might take a transaction object and return a new `Block` instance. The Factory Pattern encapsulates the instantiation logic (e.g. different block types) outside client code <sup>11</sup>.
- **DAO (Data Access Objects):** All database operations go through DAO classes. This keeps SQL/SQLite code separate from business logic <sup>7</sup>. For instance, `BlockDao` or `TransactionDao` handle `INSERT` / `SELECT` statements, while `LedgerService` does not manage SQL directly.
- **Others:** The overall architecture could follow **MVC** for the desktop GUI or web UI. Dependency Injection (manual or via frameworks) can be used to wire services/DAOs. Consistent use of logging, exception handling, and configuration management (via Maven profiles) are also best practices.

Using these patterns yields robust, maintainable code. For example, the DAO pattern “decouples the data persistence logic to a separate layer” <sup>7</sup>, which simplifies testing and future changes.

## Sample Data and Code Snippets

**Sample Block and Transaction:** For testing, you can manually create some blocks and transactions. For example, pseudo-code inspired by a typical blockchain demo <sup>12</sup>:

```
Blockchain blockchain = new Blockchain();
blockchain.addBlock(new Block("First transaction: Alice pays Bob $100",
"0")); // Genesis block
blockchain.addBlock(new Block("Second transaction: Bob receives $100",
blockchain.getChain().get(0).hash));
```

This creates a genesis block ( `previousHash = "0"` ) with data and a second block that links to the first. In SQLite, you might see rows like:

```
blocks table:
+-----+-----+-----+-----+
+-----+
| block_id | timestamp | data | previous_hash |
hash |
+-----+-----+-----+-----+
+-----+
| 1 | 1689466745871 | "First transaction: Alice pays Bob $100" |
"0" | "a3f1...9b2c" |
| 2 | 1689466745905 | "Second transaction: Bob receives $100" |
"<hash of block1>" | "c4d2...7f5e" |
+-----+-----+-----+-----+
+-----+
```

```

transactions table:
+-----+-----+-----+-----+-----+
+-----+
| txn_id | block_id | type      | amount | description |
| txn_timestamp |
+-----+-----+-----+-----+-----+
+-----+
| 101    | 1        | FINANCIAL | 100.0  | "Alice -> Bob payment" |
1689466745871 |
| 102    | 2        | FINANCIAL | 100.0  | "Bob receives payment" |
1689466745905 |
+-----+-----+-----+-----+-----+
+-----+

```

### Code Snippet (SHA-256):

```

String input = previousHash + timeStamp + data;
MessageDigest md = MessageDigest.getInstance("SHA-256");
byte[] digest = md.digest(input.getBytes(StandardCharsets.UTF_8));
String hash = bytesToHex(digest); // convert bytes to hex string

```

This pattern is standard in Java for SHA-256 <sup>13</sup>. The helper `bytesToHex` simply formats the bytes as hexadecimal.

### Sample SQLite Connection (JDBC):

```

String url = "jdbc:sqlite:/path/to/quantumledger.db";
Connection conn = DriverManager.getConnection(url);
System.out.println("Connected to SQLite database");

```

Here `org.sqlite.JDBC` is loaded automatically by the driver; this uses the SQLite file-based URL <sup>14</sup>.

These examples illustrate how blocks and transactions flow through the system. Extensive unit and integration tests should populate dummy data to verify the ledger's immutability and query results.

## References

The above design and examples follow common blockchain and Java practices. For instance, a Java blockchain tutorial shows a block class with fields for `hash`, `previousHash`, `data`, and `timeStamp` that calculates its hash using SHA-256 <sup>2</sup>. The concept of an immutable ledger is well-known: once data is added it "cannot be altered without changing all subsequent blocks" <sup>5</sup>. Patterns such as Singleton, Factory, and DAO are widely recommended for Java applications <sup>15</sup> <sup>7</sup>. SQLite is chosen as a simple embedded database: it is lightweight, reliable, and often used for desktop financial applications <sup>16</sup>. The JDBC connection string for SQLite is typically `jdbc:sqlite:path_to_db_file` <sup>14</sup>. Finally, Hibernate 6 offers a built-in SQLite dialect, easing a future migration from raw JDBC to JPA/Hibernate <sup>1</sup>. These sources and examples guided the design choices above, ensuring a robust, scalable ledger system.



- 
- 1 **java - Does Hibernate Fully Support SQLite? - Stack Overflow**  
<https://stackoverflow.com/questions/17587753/does-hibernate-fully-support-sqlite>
  - 2 3 12 **Implementation of Blockchain in Java - GeeksforGeeks**  
<https://www.geeksforgeeks.org/implementation-of-blockchain-in-java/>
  - 4 **BlockChain.java - It contains list of blocks · GitHub**  
<https://gist.github.com/madan712/463821142be0c5e3ff17f9fadcece031>
  - 5 6 **What Is Blockchain Technology? Understanding the Basics**  
<https://atomicwallet.io/academy/articles/what-is-blockchain>
  - 7 10 11 15 **Most Common Design Patterns in Java (with Examples) | DigitalOcean**  
<https://www.digitalocean.com/community/tutorials/java-design-patterns-example-tutorial>
  - 8 **Enforcing Immutability in your Double-Entry Ledger**  
<https://www.moderntreasury.com/journal/enforcing-immutability-in-your-double-entry-ledger>
  - 9 **Designing Accounting Database - General Ledger vs Sub-Ledgers - Stack Overflow**  
<https://stackoverflow.com/questions/68458637/designing-accounting-database-general-ledger-vs-sub-ledgers>
  - 13 **SHA-256 Hash in Java - GeeksforGeeks**  
<https://www.geeksforgeeks.org/java/sha-256-hash-in-java/>
  - 14 **SQLite Java: Connect to a SQLite Database using JDBC Driver**  
<https://www.sqlitetutorial.net/sqlite-java/sqlite-jdbc-driver/>
  - 16 **Appropriate Uses For SQLite**  
<https://www.sqlite.org/whentouse.html>